

RUSHI ASHOK SAVANI

+1 315 952 4519 | rushiashokbhaisavani@gmail.com | [Linkedin](#) | [Github](#) | [Portfolio](#)

EDUCATION

Syracuse University, NY, USA

Masters | Information Systems, Data Science

Aug 2024 - May 2026

(GPA: 3.86/4)

- *Relevant Coursework:* Applied Machine Learning, Database Management, Data Science, Cloud Management

Thadomal Shahani Engineering College, Mumbai, India

Bachelors | Information Technology (*Minors - Artificial Intelligence and Data Science*)

Aug 2020 - Jun 2024

(GPA: 8.52/10)

WORK EXPERIENCE

AI Engineer / Data Science Co-op

ProHealth Home Health and Hospice

Jun 2025 – May 2026

Atlanta GA

- Built a LangGraph-powered clinical assistant in Snowflake, using Cortex Search to retrieve patient data fragmented across tables, helping administrators query records in natural language, saving 20+ hours weekly.
- Ran RAGAS against a 50+ question golden set built with administrators, found low context precision and recall, then layered in hybrid retrieval, lifting accuracy by 35%.
- Rebuilt the XGBoost ML model on AWS SageMaker, pruning features via SHAP and VIF, cross validating with Stratified K Fold, then threshold tuning off the AUC-PR curve to prioritize recall, lifting recall 62% to 85% and cutting readmissions 10%.
- Architected 3 LLM driven pipelines that generate citation backed root cause reports on flagged cases clinicians lacked time to manually investigate, surfacing roughly \$20K in preventable opportunities within Q1.
- Selected models through A/B testing and an LLM as judge scored against a custom rubric with manual review and maintained prompt versions through DeepEval hallucination and factual consistency checks.
- Deployed SageMaker endpoints and the investigation pipeline behind FastAPI, Docker, EC2, and GitHub Actions CI/CD, tracking, latency, errors, and prediction drift through CloudWatch.

Data Scientist - Research Assistant

Syracuse University

Jan 2025 – May 2025

Syracuse, NY

- Automated a Python pipeline that handled missing values, standardized unstructured text and normalized mixed language responses into English before fine-tuning, saving 15+ hours.
- Fine-tuned RoBERTa in PyTorch via Hugging Face Transformers on ~7,000 survey responses, tuning optimal learning-rate through loss curve with AdamW, reaching 91% accuracy under stratified K-fold validation.
- Clustered open responses using Sentence-BERT embeddings and HDBSCAN, then exposed AI adoption themes through a Streamlit dashboard built for non-technical research stakeholders.

Data Analyst

Krishna Jewels

Jun 2023 - Jun 2024

Mumbai, IN

- Migrated 40K+ paper based inventory, sales, and supplier records into a normalized MySQL schema, cutting record lookup time by 70% through indexing and query rewrites.
- Automated 15+ hours of weekly manual analysis by building SQL ETL pipelines into Power BI with DAX measures for sales, margin, and inventory KPIs, building interactive dashboards.

TECHNICAL SKILLS

- **Languages & Data:** Python, SQL, PostgreSQL, MySQL, TypeScript, Git, REST APIs
- **LLM & GenAI:** LangGraph, RAG, Vector Databases, Prompt Engineering, DeepEval, RAGAS
- **Machine Learning:** Scikit-learn, XGBoost, Classification, Regression, MLflow, SHAP, Cross Validation, A/B Testing
- **Deep Learning & NLP:** PyTorch, Hugging Face Transformers, RoBERTa, ClinicalBERT, Sentence-BERT, TF-IDF
- **Analytics & BI:** Power BI, Tableau, Streamlit, ETL/ELT, Data Modeling, Statistical Analysis, Hypothesis Testing
- **MLOps & Cloud:** AWS SageMaker, Snowflake, Docker, FastAPI, GitHub Actions, CI/CD, CloudWatch, Model Monitor

PROJECTS

PrivateRAG | *Fully Local Private-Documents Q&A System*, [GitHub](#)

- Designed a fully local Retrieval-Augmented Generation (RAG) stack with document ingestion, Qdrant vector indexing, metadata filtering, and Llama 3.1 inference with zero external API calls.
- Wrote a TypeScript chat interface over a local OpenAI-compatible backend with streaming responses, source attribution, and chunk-level traceability.

COVID-19 India Analytics | *Interactive Tableau Dashboard*, [Dashboard](#)

- Visualized India's COVID-19 case, testing, and vaccination trends in one interactive Tableau dashboard.
- Wrangled several large public datasets in Python, reconciling inconsistent state-level records before loading them into Tableau for analysis.
- Mapped vaccine dose administration against testing capacity, exposing regional disparities across age groups

KarmaScore Predictor | *Subscription Churn Analysis*

- Trained an XGBoost churn classifier on 30K+ subscriber records, narrowing 60 candidate variables to 23 predictors through recursive feature elimination and mutual information scoring.
- Ranked subscribers into decile risk bands using PR-AUC and lift analysis, producing a prioritized list for targeted retention planning.